PRODUCT OVERVIEW - About SPOCK - AI Native Discovery Tool For Product Management Lifecycle

1. What is SPOCK?

Overall definition

SPOCK is eská spoitelna's internal agentic product for product discovery. It is intended to become the environment in which product teams do structured product work — not a general-purpose AI chat tool added on top of the current process.

The operating principle holds across both the pilot and the full vision: **the agents prepare the work, people decide.** Every approval, selection and sign-off is made and recorded by a named person with their role; no agent approves, selects or signs anything on its own.

SPOCK today, in its pilot phase

In its current pilot form, SPOCK runs as one connected agentic workflow inside GitHub Copilot in VS Code. It takes an initiative from a brief to a delivery-ready initiative document in **seven stages**, carried by **four workflow agents**, with **two standalone agents** a person can call at any time. The method behind the workflow is the Opportunity Solution Tree: a brief commits to an outcome, problems are found in material the team already holds, they are framed as opportunities, solutions are bound to the opportunities they serve, and everything after a chosen solution narrows it towards delivery.

SPOCK's imagined full scope

This is the future-state vision, not the current pilot scope. The intended end-to-end flow runs from an initial idea or problem-space brief, through research (market, competitive, trend and customer), through opportunity framing (desired outcome, opportunity definition, solution brainstorming, prioritization), into roadmapping (first product increment, longer-term direction) and prototype, into product specifications (functional and non-functional analysis), and on to a connection with delivery. Other approved systems and agents — including the engineering AI delivery agents being developed separately — may connect to SPOCK in future through a controlled, governed approach.

The current pilot covers only the front part of this flow: it takes a brief through market research and onto a single first increment (MVP) and a delivery-ready initiative document. Multi-increment roadmapping and formal functional/non-functional specification as well as customer research and prototyping are all descoped from the pilot and remain future-state capabilities.

2. What problem are we solving?

The key problems in how product teams work today that SPOCK is intended to address:

- **Inconsistent discovery.** Product teams currently use different methods, workflows, research practices and documentation standards, making it hard to set a common bar for how initiatives are explored, validated, prioritized and prepared for delivery. Many teams skip discovery overall and start with unverfied solution that is scoped, analyzed and delivered.
- **Recreated work, low reuse.** Teams frequently rebuild context, research, analysis and decision rationale from scratch. Relevant work may already exist elsewhere but is hard to find, access or reuse — part of the organization has no documented record of its product work today.
- **Premature commitment.** Teams often commit to a direction before the underlying assumptions are validated, so problems with value, cost, feasibility, customer need, market fit or regulatory requirements surface late.
- **Late involvement of control functions.** Risk, compliance, security and other control functions are typically brought in close to release rather than during discovery, creating avoidable friction, rework, delay or scope change.
- **Uneven quality.** The quality of discovery depends heavily on the experience of whoever is running it, and knowledge one team generates is not consistently available to other teams.
- **Discovery often takes too long or costs too much (in people time or other costs).** Some teams dedicate one or more quarters to discovery because they rely on external consultancies or because process itself is unexpectedly complex and convoluted across mutliple teams.

3. Who is it for?

Primary users are the people running product discovery inside product teams: product managers and product designers, but also software engineers, analysts, data and other domain experts. Any PDE team members who build the brief, do research, create and select the solution and shape the MVP, define product requrements and Jira initiatives and other product / business documentation are SPOCK users.

Secondary audiences read and act on SPOCK's outputs without running the workflow themselves: business and engineering stakeholders who review solutions and the MVP; compliance, legal, GDPR/privacy and ICT-security reviewers, who receive the regulatory touchpoints each solution is screened against; and executive, governance and AI-working-group stakeholders tracking the pilot and its results.

SPOCK pilot is limited to only select number of teams from different tribes and also specifically to few select people in these teams. Plan for full rollout and scaling of SPOCK will be done after the pilot.

4. Why does the bank need it?

Discovery time and expertise are not evenly available across teams. Product discovery competes with delivery pressure in every team, and few teams have the dedicated time or seasoned product/research expertise to run it well on every initiative. The bank's own PDE talent mapping — 847 roles mapped so far across product, engineering and delivery crafts — shows a workforce that PDE teams lack seniority when it comes to strategic product management and outcome driven thinking, which is exactly why discovery quality today depends so heavily on whoever happens to be running it rather than on a shared, repeatable method. SPOCK is meant to close that gap directly: it carries the discovery method itself (the Opportunity Solution Tree, applied through gated stages) so a team can learn and adopt new practices more effectively — the workflow supplies the structure, the team supplies the judgement.

Teams don't share what they learn, so the same work gets repeated. Nothing today connects one team's market research, competitive findings, or discovery decisions to another team's. When two teams face similar problems, neither knows what the other already found — so both spend the same weeks rebuilding the same context from scratch. This is the same problem behind the product work that has no documented record at all: even when a team's discovery is good, it's invisible to everyone else. SPOCK's answer to this is only partly built today — Pilot Definition itself keeps each run local to one team but full scope SPOCK will look to solve for cross-team knowledge and its reuse.

The pace of technology change makes this urgent, not optional. Software delivery speed across the industry is being reset by AI-assisted tooling — teams that adopt AI-native ways of working are moving markedly faster than teams that don't, and that gap is only widening. eská spoitelna is developing its own AI native ways of working across areas and teams and SPOCK is currently one of few lighthouse projects, alongside agentic software engineering, working out how the bank provisions AI capability at all.

5. What outcomes are we targeting?

For product teams

- Better product decissions based in data and made using well proven product decission-making frameworks.
- Less time producing first drafts, more time on judgement and decisions.
- A consistent, traceable record of who decided what, when and why, built into every stage.
- Solutions that already carry their regulatory touchpoints instead of discovering them late.
- A faster product discovery path from a supplied brief to a delivery-ready initiative document.

For eská spoitelna

- A shared discovery standard usable across multiple product teams, rather than one that depends on who happens to be running it.
- A more effective vechicle to teach & coach PDE teams and any future new team members on ‘how we build products in eská spoitelna’ – workshops and hands on support remain the work of CoEs, but are reserved for high impact work, while working with SPOCK and consulting its coaching agents covers the basics.
- Earlier visibility of risk, compliance, security, legal and GDPR/privacy considerations.
- A documented decision trail for initiatives entering delivery.
- A working, governed example of AI adoption inside a regulated environment that the wider bank-wide AI strategy can build on.

6. How will we measure success in the long term?

The measures below are for tracking SPOCK once it runs beyond pilot and a single testing event. They are proposed here based on how the market currently measures long-running internal AI and agent tools. None of the targets below have been set yet.

Metric	What it measures	Why it matters	Status today
SPOCK-origin ratio	Share of initiatives entering delivery that went through a SPOCK-run discovery, versus discovery done outside the workflow.	Distinguishes real reliance on the tool from occasional or token use — usage alone does not show how much real organizational output actually flows through it.	MISSING: not yet tracked — requires linking SPOCK runs to delivery-system records
Trust index	Share of SPOCK outputs accepted without material rewrite, tracked continuously by stage over time, not measured once.	Trust does not move with adoption and can decline even as usage rises; a single reading would miss that trend.	Measurable in principle now — approval without rewrite is already a recorded decision — but only meaningful once tracked as a trend across many runs
Downstream durability / rework rate	Share of SPOCK-originated decisions (opportunities, selected solution, MVP scope) that reach delivery unchanged, versus get reworked or reversed after handoff.	Confirms discovery speed is a genuine gain rather than work resurfacing later as delivery rework.	MISSING: not yet tracked — requires a link between SPOCK's output and delivery outcomes
Time-to-validated idea **Lead time** **Time-to-value**	Time from brief acceptance to each major gate, tracked as a trend across every initiative run through SPOCK over time, not a single measurement.	Shows whether time savings hold, grow or fade as usage scales past the first pilot teams.	Measurable in principle now as part of PDE metrics tracked on every initiative. But needs linking to SPOCK

7. What is the relationship between humans and SPOCK agents?

Principle: the agents prepare the work, people decide. No agent approves, selects or signs anything, and no agent judges whether its own output passed.

What agents do

- Draft and structure each stage's output from the material and decisions available to them.
- Run an automated quality gate on every stage's output for completeness, form and traceability, before a person sees it.
- Screen each solution against four regulatory areas and carry the touchpoints forward into the MVP and PRD.

What humans do

- Supply the problem material the workflow starts from.
- Make every substantive decision in the workflow — which opportunities to pursue, which solution to select, which MVP option to build, whether the PRD and the initiative are ready.
- Record each decision under their own name and role.
- Check outputs for accuracy and completeness and challenge unsupported statements.
- Sign any exception where a stage is allowed to proceed despite a failed gate.

Where humans decide

Gate	Stage	Who	What they decide
G1	Brief	Product Manager	The brief is complete enough to proceed
G2	Problems	Researcher and Product Manager	Each opportunity is well formed and evidenced, and which ones the run pursues
—	Market	nobody	Working material only; no sign-off — read inside the Solutions stage
G4	Solutions	Product Manager	One solution selected
—	MVP	Product Manager	One MVP option chosen
G5b	MVP	Product Manager	The expanded MVP accepted
G6	PRD	Product Manager	The PRD agreed
G7	Initiative	Product Manager	The initiative approved as a delivery item

8. What agents cannot do autonomously

- Approve, select or sign off anything.
- Judge whether their own output is correct.
- Make the compliance judgement — the workflow screens and lists touchpoints, but whether a solution is compliant stays with the compliance teams.
- Create anything in Jira, or estimate effort, cost or timeline.

9. Who owns and develops SPOCK?

SPOCK is developed and governed centrally by the CoE Product Management and CoE Product Design who own the agentic product, its methodology, workflows, distribution / scalling and roadmap. The quality of individual capabilities is reviewed and improved by the relevant experts closest to that work, mostly Guild masters in both CoEs who are accountable for quality of outputs of SPOCK and its usefulness for PDE teams across eská spoitelna tribes.

Within the bank's still-forming new AI native ways of work SPOCK is currently one of few lighthouse initiatives — alongside agentic software engineering — both used to work out how the bank provisions and governs agents that make product and engineering teams more effective and advanced in their everyday work.

For more info contact:

Bambura-Alibaši Ines ibamburac-alibasic@csas.cz

Bedáová Veronika vbedanova@csas.cz